

Designing a Virtual Security Management Layer for Cloud Content

Lukas Klingsbo

January 28, 2016

Abstract

TODO: Abstract

Keywords:

Contents

1	Introduction	1
1.1	Prior Work	1
1.1.1	Copy-on-Write	1
2	Related Terminology	2
2.1	Abbreviations	2
2.1.1	JPF	2
2.1.2	CDN	2
2.2	Terms	2
2.2.1	Snapshot	2
3	Model	3
3.1	Related work	3
3.2	Approach	3
3.3	Elements of the Model	3
3.4	Access rights	3
3.5	Data integrity	4
3.6	Initial Assumptions	4
3.7	Integrity Threats	4
3.8	Findings	5
4	Implementation	6
4.1	Background	6
4.1.1	About Uprise	6
4.1.2	The current system	6
4.1.3	Problem description	6
4.2	Related Technologies	6
4.2.1	React	6
4.2.2	Reflux	7
4.2.3	Scala	7
4.2.4	REST	7
4.2.5	MongoDB	7
4.3	Methods for determining implementation details	8
4.4	Snapshot functionality	8
4.4.1	Copy-on-Write	8
4.4.2	Full Copy	8
4.4.3	Comparison of Copy-on-Write system implementations	8

4.4.4	Snapshot functionality of Perius	9
4.5	Resulting system	10
4.5.1	Perius	10
4.5.2	Copy-on-Write	11
4.5.3	Persistent storage	11
4.5.4	API	11
4.6	Findings	13
4.6.1	Scalability	13
4.6.2	Security of the system	16
5	JPF	16
5.1	Entities	17
5.1.1	Content	17
5.1.2	Project	17
5.1.3	Container	17
5.1.4	Snapshot	17
5.1.5	File	18
5.1.6	User	18
5.2	Execution	18
6	Discussion	19
7	Summary	19
7.1	Conclusions	19
7.2	Future work	19

1 Introduction

Developing larger projects containing static content usually involves using a Content Distribution Network to be able to scale to a large user base. The commercial Content Distribution Networks are usually fairly easy to use, the content that is to be used in a project is usually simply uploaded and then distributed over the globe when the public requests it. For secret content this can be a problem and an inconvenience, and that is what this thesis is about. This work examines ways of enforcing virtual access control on content and groups of content, in the form of folders and snapshots. A system was developed to make the underlying theory work in practice.

The research question that this report answers is how and whether it is practically feasible to use Copy-on-Write for a high-level system like the one that is implemented.

TODO: Clarify problem definition and research questions

1.1 Prior Work

1.1.1 Copy-on-Write

This work relies heavily on the Copy-on-Write principle, which was founded and used in the Mach kernel [1], as it can be used to efficiently create snapshots and help solving concurrency problems that otherwise can occur.

Copy-on-Write is used in for example virtual memory management systems [2], for snapshot and as an optimisation technique for objects and types in several programming languages [3].

Its principle is that when processes or nodes share data in between each other, the data is not copied until one of the processes does changes to it. This is an optimisation as the processes does not have to send or copy all of the related data that is in memory, rather they only have to send pointers to the data. After many Copy-on-Write's a complex tree structure can be built up, but optimisations can be done to simplify that structure [4].

TODO: Polish and extend

2 Related Terminology

2.1 Abbreviations

2.1.1 JPF

Java Path Finder - It was developed by NASA and in 2005 they released it under an open source licence, which made more people contribute to the project. JPF is usually used for doing model checking of concurrent programs to easily find for example race conditions and dead locks.

2.1.2 CDN

Content Distribution/Delivery Network - Replicates content to several servers, usually spread out geographically. Once a request is made, the network serves content from the server closest to the requester.

2.2 Terms

2.2.1 Snapshot

A snapshot is a way to record the full state of a system at a specific time. The term comes from photography where a photo can be seen as the state of what the photo is of, at a certain time. Snapshots should not be confused with full copies of a system as full copies can be used as backups meanwhile snapshots are not very effective means of backups in the case of data corruption. It is not effective against data corruption as snapshots usually still refer to unchanged data that is still a part of the system [5].

3 Model

The model for this work should show how the data can not be accessed or modified by unauthorized users and how the integrity of the data is always kept in the Perius system.

There could also be another relevant model done to show that content can not be accessed by unauthorized viewers once the content is uploaded to a CDN, but as that should already have been thoroughly checked by the CDN providers this work can focus solely on the internal users and content of the management system.

3.1 Related work

3.2 Approach

3.3 Elements of the Model

TODO: Mathematical formal description of how the system works

Set	Elements	semantics
C	$c_0 \dots c_n$	Containers; folders in the virtual file system
F	$f_0 \dots f_n$	Files; files, images, videos
P	$p_0 \dots p_n$	Content; Meta-data for files
U	$u_0 \dots u_n$	Users; registered users in the system
A	$A[u_0, c_0] \dots A[u_n, c_n]$	Access matrix; describes what containers users have access to

3.4 Access rights

The a notation is meant to symbolise an access token, and if a exists in the requested entity of the matrix the corresponding user to that entity has full access to the corresponding object.

$$\begin{aligned}
 u \in U \text{ can read } c \in C &\Leftrightarrow u \in A[u, c] \\
 u \in U \text{ can write } c \in C &\Leftrightarrow u \in A[u, c] \text{ and readonly } \notin c \\
 u \in U \text{ can delete } p \in c &\Leftrightarrow u \in A[u, c] \text{ and readonly } \notin c \\
 u \notin U \text{ can delete } f \in F & \\
 \forall c \in C, \quad \exists u \in U \mid a \in A[u, c] &
 \end{aligned} \tag{1}$$

3.5 Data integrity

A computer system or subsystem is defined as possessing the property of integrity if it behaves consistently according to a defined standard. This implies that a subsystem possessing the property of integrity does not guarantee an absolute behaviour of the system, but rather that it performs according to what its creator intended [6].

3.6 Initial Assumptions

To create an integrity model, some initial assumptions have to be made about what the correct behaviour of the system is, which the model then can be shown to follow. In this work unintentional behaviour as the result of data modification is the main concern, which could be used for sabotage or simply be the effect unintentional unfortunate race conditions etc.

3.7 Integrity Threats

According to Biba et. al [6] one can consider two threat sources, namely subsystem external and subsystem internal. The external sources could be another system calling the subsystem with faulty data or trying to make inaccurate calls to program functions, it could also be somebody trying to tamper with the exposed functions of the program. Threats that are internal could be a malicious part of the subsystem or simply an incorrect part of the subsystem, which does not behave according to specification.

In this work external threats are handled as threats that can occur from what has been exposed by the API (See Section 4.5.4 and internal threats as incorrect implementation. As the server and its system are assumed safe malicious subsystems are not considered.

TODO: Translate to mathematical expressions

- if a user wants to update a file in a content, the file is copied and the original is intact
- if a user reads content and then writes to it, the content is directly changed
- if a user copies content, a new content is created at the destination with reference to the same file
- if a user reads from a container and then writes to it, the container is directly changed (Which means last write wins, which doesn't matter)
- if a user creates a snapshot of a container, the full container tree is re created with new ids at the destination, its content still refers to the same files.

3.8 Findings

4 Implementation

4.1 Background

4.1.1 About Uprise

Uprise is a company based in Uppsala, Sweden. It is an EA studio focussing on creating great gaming experiences, which means that they are mostly not focussed on the actual gameplay which other EA studios like DICE is.

At the moment they are for example very involved with producing the Star Wars Battlefront game by making menu systems and developing the companion app, Battlefront companion.

4.1.2 The current system

Today a system called battlebinary [7] is used for managing and uploading files, mostly images, to content delivery networks. The current system does not make use out of the security features that the CDN's are offering, instead it uses a form of security by obscurity. When a file is uploaded to a CDN it is open for the public, but its filename is composed out of its original filename concatenated with a part of the MD5 hash of the content of the file, which makes it an extremely hard process to access the file on the CDN without access to the original file or a reference to the URI.

In the current system you can only upload a file once as there will be a collision in the upload otherwise, as the old and the new file will have the same MD5 hash.

4.1.3 Problem description

As the current system does not offer proper security measurements, is lacking a lot of features that is needed and does not scale very well, a new system should be developed. This work is about examining a way of implementing Copy-on-Write in a high level system like this, which should solve the scalability problem and make it possible to implement wanted features like snapshots, cloning and concurrent modifications of content.

4.2 Related Technologies

4.2.1 React

React is a JavaScript library for building user interfaces. React uses both its own virtual DOM and the browser's, this makes it able to efficiently update dynamic web pages after a change of state through comparing the old virtual DOM with the

resulting virtual DOM after the state change and then only update the browser's DOM according to the delta between the virtual DOMs [8]. React can be seen as the system for handling views in front-ends implementing a MVC (Model-View-Controller) architecture.

4.2.2 Reflux

Reflux [9] is an idea and a simple library of how to structure your application. It features a unidirectional dataflow (see Figure 1) which makes it more suitable, than for example Flux [10], when using a functional reactive programming style.




Figure 1: Reflux unidirectional dataflow

4.2.3 Scala

Scala is a multi-paradigm programming language. It most commonly runs on the JVM and compared to Java it supports most functional programming features at the same time as it supports object oriented programming [11].

4.2.4 REST

REST stands for representational state transfer, it is an architectural idea for writing stateless services. These services usually use URIs to identify specific resources and HTTP to modify or query these resources [12].

4.2.5 MongoDB

MongoDB is a document-oriented database which means that it does not have the concept of rows as normal relational databases has. Instead each entity in the database is stored as a document which is not fixed to a predefined table structure [13]. MongoDB lacks the support for joins to improve its possibility to scale, which can be a big down side to some applications containing the need for such logic.

4.3 Methods for determining implementation details

This chapter introduces the different methods used to determine how the new system should be implemented, which DBMS it should use and how the estimation of long term scaling was done.

4.4 Snapshot functionality

TODO: Structure to compare snapshot systems and conclude how Perius snapshot system was designed

4.4.1 Copy-on-Write

To efficiently create snapshots of a system Copy-on-Write can be used to make it possible to create snapshots in $O(1)$ [14], this is due to the fact that to create a snapshot in a system using Copy-on-Write you only need to reference the current nodes in the tree and make sure that they are not removed, see Figure 2.

As the persistent storage, used in this implementation (Section 4.5.3), does not implement transactions or locks a lot of different problems can occur when several clients are working on the same data set at the same time. Such problems could be race conditions and determining the happened-before relation. In this work this problem is solved by implementing Copy-on-Write. **TODO:** Move last paragraph

4.4.2 Full Copy

Full copy or deep copy, as opposed to copy-on-write, is a copy where everything is copied directly and not only when an object is changed. This is easier to implement but is in most cases more inefficient as more disk space will have to be used and if used with for example certain tree structures the part of the tree that needs to be copied will have to be traversed.

4.4.3 Comparison of Copy-on-Write system implementations

4.4.3.1 BTRFS

Btrfs is a B-tree file system for Linux which makes use of Copy-on-Write to make it able to do efficient writeable snapshots and clones. It also supports cloning of subtrees without having to actually copy the whole subtree, this is due to the Copy-on-Write effect. As several nodes in the tree can refer to the same node each node keeps track of how many parents it has by a reference counter so that the node can be deallocated once the node does not have any parents any more. The

reference counter is not stored in the nodes themselves but rather in a separate data structure so that a nodes counter can be modified without modifying the node itself and therefore eludes the Copy-on-Write that would have to occur.




Figure 2: Cloning mechanism of Btrfs [14]

4.4.3.2 Mach kernel

In the mid 80's when the development of the Mach kernel started, there was problems with that physically copying memory was too slow. Too minimise the copying of memory, Copy-on-Write was implemented. It was implemented so that virtual copy operations could be done and so that tasks could share read-write memory [15].

TODO: Insert more systems here

4.4.4 Snapshot functionality of Perius

In Perius snapshots and clones are not taken in the fashion which Btrfs uses, which can be seen in Figure 2. As Perius does not have the tree structure pre-built and each node is instead stored in a flat storage space, such operation would be too computationally expensive as trees would have to be merged when collisions occur, due to the non-blocking nature of the application. Instead this implementation makes a full copy of the meta-data of the tree, but still refers to the same binary files until they are changed, which results in the creation of a new node.

TODO: Relate section more to comparison

4.5 Resulting system

4.5.1 Perius

Perius is the implementation that was done to solve the problem at hand at Uprise. Perius has a back-end written in Scala and a front-end written in Javascript (ES6), but they are both interchangeable. The back-end has a REST API running, which is how the front-end communicates with the back-end.

TODO: Picture of the newest front-end




Figure 3: Front-end [14]

The service features a virtual file structure over the assets that has been stored, snapshots, security management of whole containers as well as individual files, audit and access logging, multi project support and a modular design for persistent storage.

The front-end is written in ES6 with React and Reflux, and the styling is done with the help of Bootstrap.

4.5.2 Copy-on-Write

TODO: Rename section and move

This implementation is far from as efficient as the other Copy-on-Write systems described in Section 4.4.1 in most aspects, but more efficient in some. As the implementation is built upon MongoDB as persistent storage and not a pure tree structure, single nodes can be fetched in $O(1)$ but when querying for subtrees they need to be built first, which takes $O(\log(n))$, where n is the number of nodes in the subtree.

4.5.3 Persistent storage

4.5.3.1 MongoDB

MongoDB was chosen as the persistent storage because of its quick lookups and because of its internal storage format called BSON, which is very similar to JSON which the API is using. As the formats are similar, the process of marshalling and unmarshalling becomes quite easy between the core code, MongoDB instance and REST interface. The second reason was that if the system needs to scale in the future it is very easy to distribute MongoDB and if needed the system can easily be migrated to Reactive Mongo, which is an asynchronous and non-blocking driver for MongoDB and can therefore make the system scale even further [16].

All files are also stored directly in MongoDB with the help of GridFS. GridFS chunks the files according to the size limit of MongoDB objects, which is currently 4MB. The advantage of this is that backups of the Perius state is easily done through a database backup, no separate files needs to be backed up. Another advantage that is given by this is that you can retrieve specific ranges of a file, although that advantage is not needed in the Perius implementation. The disadvantage of using this approach is that when using a non-distributed database it is slower to read and write to the database than reading or writing directly to the filesystem. Another disadvantage is that to access the files it is needed to go through the database layer in some way, instead of accessing the filesystem.

4.5.4 API

In this work a RESTful API was implemented and used for back-end \leftrightarrow front-end communication.

REST was chosen as only basic CRUD operations needs to be performed and because the BSON format which is used in MongoDB is almost identical [17] to the standardised JSON format which is usually used by RESTful services [18].

4.5.4.1 REST Endpoints

For the front-end to communicate with the back-end, a RESTful service is implemented. The following endpoints were configured:

- projects
 - GET - list all projects
 - POST - create new project
- projects/{id}
 - GET - get specific project
 - PUT - update existing project
 - DELETE - delete existing project
- projects/{id}/content
 - POST - create new content in a specific project
- projects/{id}/content/{id}
 - GET - get specific content in a specific project
 - PUT - update existing content in a specific project
 - DELETE - delete existing content in a specific project
- projects/{id}/snapshots
 - POST - create new snapshot in a specific project
- projects/{id}/containers
 - POST - create new container in a specific project
- projects/{id}/containers/{id}
 - GET - get specific container in a specific project
 - PUT - update existing container
 - DELETE - delete existing container

As can be seen several expected endpoints are missing, this is intentional as the operations missing can be performed in a more efficient way. Such endpoint is for example *GET projects/{id}/containers* as all containers exist in *GET projects/{id}* and the interface should present a file structure where both content and containers are shown.

4.6 Findings

4.6.1 Scalability

When using the ReactiveMongo driver [16], which is asynchronous and non-blocking, the application has no limits of how much load and users it can handle as the hardware and nodes can be scaled up linearly when needed. With Cashbah [19], which is used with the current implementation, it is harder to scale to the enormous amounts of load which ReactiveMongo can support as Cashbah is synchronous and has blocking IO. For this work the kind of scalability which is offered by ReactiveMongo is not needed as the load will not reach the peak (**TODO**: relate to graph) for what Cashbah can handle on a single server.

4.6.1.1 Load Testing

Wrk was used to load test the back-end. It is a multi-threaded benchmarking tool for HTTP which can create large loads. The testing was done locally on a virtual server having 6 CPU cores and 48GB of RAM. The results that will be analysed are latency and throughput of the server as that will give a fair judgement of how well the system can perform.

GET requests To be able to determine what is the bottleneck of the back-end three different types of *GET requests* are performed. The first one tries to maximise the number of requests that spray-can (HTTP server) can handle with the given specifications and therefore the server just returns 200 (OK). The second request requests the simplest API call which involves getting all currently stored projects from the MongoDB instance. As this involves the server answering with some payload, the third request is simply the same amount of static payload as the second request but without involving any connections to MongoDB.

The typical requests with Wrk in this work will look similar to this:

```
wrk -t100 -c100 -d10s http://perius:8000/ok
```

where -t100 means that it uses 100 threads, -c100 that it emulates 100 clients requesting over and over and -d10s is the amount of time that the load test will run, in this case 100 seconds.

Listing 1: Result of OK requests

```
wrk -t100 -c100 -d10s http://perius:8000/ok

Running 2m test @ http://perius:8000/ok
100 threads and 100 connections
  Thread Stats   Avg      Stdev     Max   +/-  Stdev
    Latency    1.22ms    3.20ms 166.78ms   97.75%
    Req/Sec    1.04k    152.50   2.25k    79.18%
 10345746 requests in 1.67m, 1.69GB read
Requests/sec: 103353.85
Transfer/sec:   17.25MB
```

Listing 2: Result of MongoDB requests

```
wrk -t100 -c100 -d10s http://perius:8000/projects

Running 2m test @ http://perius:8000/projects
100 threads and 100 connections
  Thread Stats   Avg      Stdev     Max   +/-  Stdev
    Latency   14.52ms   13.34ms 375.51ms   99.31%
    Req/Sec    72.85     6.91   232.00    60.20%
 726351 requests in 1.67m, 196.03MB read
Requests/sec: 7256.38
Transfer/sec:   1.96MB
```

Listing 3: Result of static text requests

```
wrk -t100 -c100 -d10s http://perius:8000/static

Running 2m test @ http://perius:8000/static
100 threads and 100 connections
  Thread Stats   Avg      Stdev     Max   +/-  Stdev
    Latency    6.42ms   28.99ms 290.37ms   96.26%
    Req/Sec    0.94k    203.30   2.00k    83.83%
 9199737 requests in 1.67m, 2.19GB read
Requests/sec: 91902.48
Transfer/sec:   22.44MB
```

Listing 1 shows that the HTTP server can answer about 100K requests/s when not involving any payload other than status code 200 (OK). When comparing that to the request which involved the server responding with a small (100 Byte)

payload (Listing 3 one can see that it is about 10% slower to send some more data, ~90K requests/s. However when comparing that to the requests that needed database access it is obvious that the HTTP server can handle all the load that it needs to, as Listing 2 shows that to fetch all documents in a collection (in this case only one) is dramatically slower, the throughput was only ~7000 requests/s.

These are very simple requests, to evaluate how much load the application can handle in its current state a more complex but common request has to be analysed. The most complicated and still common request that system is receiving is requests of a full project tree. This request is computationally expensive as the tree is not stored directly in the database but has to be built from the id and parent id of each container and content, a simulated project tree of similar size to the ones stored in the current management system reaches around 25KB in uncompressed size.

Listing 4: Result of project tree requests						
wrk -t10 -c10 -d10s http://perius:8000/projects/56a5f...						
Running 10s test @ http://perius:8000/projects/56a5f...						
10 threads and 10 connections						
Thread Stats	Avg	Stdev	Max	+/-	Stdev	
Latency	247.54ms	21.64ms	323.78ms	96.76%		
Req/Sec	3.80	1.22	20.00	99.00%		
401 requests in 10.07s, 9.67MB read						
Requests/sec:	39.82					
Transfer/sec:	0.96MB					

The result in Listing 4 shows an average of 40 requests/s, which means this is the real bottleneck in the application. There are two solutions to fix this bottleneck, as the project tree wont be modified nearly as often as it is requested the first solution would be to cache the results of the tree requests and invalidate those caches when the tree is modified. The second solution would be to make a more computationally feasible way of building and fetching the whole tree from the database. This is left to implement if the need comes for it in the future. Even though this massive bottleneck is present, it wont be a problem in a small scale production environment as the tree is fetched approximately once every 30 seconds by active users, which means that the application still could support at least 1200 very active users.

Listing 5: Result of cached project tree requests	
wrk -t100 -c100 -d100s	
http://perius:8000/projects/cached/56a5f...	

```

Running 2m test @
http://perius:8000/projects/cached/56a5f...
100 threads and 100 connections
Thread Stats      Avg          Stdev         Max    +/-  Stdev
  Latency        1.46ms        3.62ms    180.68ms    97.87%
  Req/Sec    838.58        120.95      1.74k     79.12%
8343565 requests in 1.67m, 17.03GB read
Requests/sec: 83352.77
Transfer/sec: 174.17MB

```

With caching turned on, on a single back-end, the application can theoretically support several million active users which are not posting any content, this is simulated in Listing 5.

POST requests

TODO: Add graphs of load testing etc

4.6.2 Security of the system

Authorization The authorization of users in the system is currently being handled by LDAP [20], as Perius is mainly focussed at being deployed in internal networks which usually has an LDAP service enabled. LDAP also makes it easier for the user to login as no separate account is needed for Perius and thus the user can use the same account as for all the LDAP connected services on the internal network.

Audit logs

5 JPF

TODO: Not sure where to put this section

JPF was used to test the idea and state transitioning of the application. Its a very simplified version of the real system that still contains all the important Copy-on-Write core concepts and the assumptions that have been made for the model. This simplified version could then be automatically tested for soundness. It is not a proof that the model works, but it is very exhaustive in its testing.




Figure 4: High Level Entity Relationships

5.1 Entities

5.1.1 Content

Content is meta data about a file and is stored in a container, it is a form of virtual file. The content can refer to for example an image, video or binary blob.

5.1.2 Project

A project is what is created to contain all content and containers related to a real project. Files can be changed within a project and the system can contain several projects and their virtual content are completely disjoint.

5.1.3 Container

A container is a virtual folder within a project which can contain content and other containers.

5.1.4 Snapshot

A snapshot is a read-only container from the state which the container the was in when the snapshot was created. A snapshot can not be updated and can only be deleted from the root of the snapshot. Snapshots are by default stored as siblings to the container which they were made from, but they can be contained by any container.

5.1.5 File

A file refers to an actual physical file. Files are stored in the database to make backup, deployment and migration easier.

5.1.6 User

A user is the structure that handles people who have been granted access to the system. Access to the system is handled by a separate service, like LDAP.

5.2 Execution

Java path finder was used to show that the model and plan of how to build the system was sound. The model was built in Java with the objective of being as reduced and simple as possible, without losing any of the cases that needed to be covered by the model checker. As the users are mainly going to be handled by external systems they were not included in the model.

Each collection in the persistent storage was emulated by using the built-in `ConcurrentHashMap` type. Each client was represented by a thread and each action taken by the client was randomised. The id hashes which MongoDB is using for each entity was imported from the `mongo-java-driver-2.13.3` and each object had its own id, generated in the same fashion as the real implementation is using, randomly generated by the `ObjectId` class to minimise collisions that is. Furthermore no locking or transactions were used and the threads were running fully concurrently, without any sleep statements.

`ConcurrentHashMap` had to be used instead of the normal `HashMap`, as the normal `HashMap`s can't be iterated over concurrently.

JPF checked each permutation of states that the threads can end up in, the result of the run can be seen in Listing 6.

Listing 6: Results of JPF run	
elapsed time:	14:26:53
states:	new=160853259, visited=451102505, backtracked=611955764, end=21640
search:	maxDepth=380, constraints=0
choice generators:	thread=160853255 (signal=0, lock=3603938, sharedRef=146989208,

	threadApi=3, reschedule=10260106), data=0
heap:	new=676056850, released=435060996, maxLive=655, gcCycles=523950061
instructions:	11917045758
max memory:	6256MB
loaded code:	classes=111, methods=2179

6 Discussion

7 Summary

7.1 Conclusions

7.2 Future work

Full access control was not implemented according to the model described in 2.2.1, it was only implemented to check whether a user should have access to the system as a whole or not, the implementation did not set or check any specific access rights to certain contents or containers.

References

- [1] M. Accetta, R. Baron, W. Bolosky, D. Golub, R. Rashid, A. Tevanian, and M. Young, “Mach: A new kernel foundation for unix development,” 1986.
- [2] J. M. Smith and G. Q. Maguire Jr, “Effects of copy-on-write memory management on the response time of unix fork operations,” *Computing Systems*, vol. 1, no. 3, pp. 255–278, 1988.
- [3] R. G. White, “Copy-on-write objects for c++,” *The C Users Journal*, 1991.
- [4] F. J. T. Fábrega, F. Javier, and J. D. Guttman, “Copy on write,” 1995.
- [5] J. Guthrie, “Method and system for taking a data snapshot,” July 8 2003. US Patent App. 10/616,411.
- [6] K. J. Biba, “Integrity considerations for secure computer systems,” tech. rep., DTIC Document, 1977.
- [7] “Github - battlelog/battlebinary,” Uprise, September 2015. http://curl.haxx.se/docs/faq.html#What_is_cURL, accessed 2015-10-15.
- [8] A. React, “Javascript library for building user interfaces,” 2014.
- [9] “Github - refluxjs,” Reflux, October 2015. <https://github.com/reflux/refluxjs>, accessed 2015-11-24.
- [10] C. Gackenheim, “Introducing flux: An application architecture for react,” in *Introduction to React*, pp. 87–106, Springer, 2015.
- [11] M. Odersky, P. Altherr, V. Cremet, B. Emir, S. Micheloud, N. Mihaylov, M. Schinz, E. Stenman, and M. Zenger, “The scala language specification,” 2004.
- [12] L. Richardson and S. Ruby, *RESTful web services*. " O'Reilly Media, Inc.", 2008.
- [13] K. Chodorow, *MongoDB: the definitive guide*. " O'Reilly Media, Inc.", 2013.
- [14] O. Rodeh, J. Bacik, and C. Mason, “Btrfs: The linux b-tree filesystem,” *ACM Transactions on Storage (TOS)*, vol. 9, no. 3, p. 9, 2013.
- [15] M. Accetta, R. Baron, W. Bolosky, D. Golub, R. Rashid, A. Tevanian, and M. Young, “Mach: A new kernel foundation for unix development,” 1986.

- [16] R. Haddock, “Intelligent internet system with adaptive user interface providing one-step access to knowledge,” Mar. 14 2014. US Patent App. 14/212,654.
- [17] D. Tomaszuk, “Document-oriented triple store based on rdf/json,” *Studies in Logic, Grammar and Rhetoric*,(22 (35)), p. 130, 2010.
- [18] L. Richardson and S. Ruby, *RESTful web services*. " O'Reilly Media, Inc.", 2008.
- [19] T. Alexandre, *Scala for Java Developers*. Packt, 2014.
- [20] T. A. Howes, M. C. Smith, and G. S. Good, *Understanding and deploying LDAP directory services*. Addison-Wesley Longman Publishing Co., Inc., 2003.